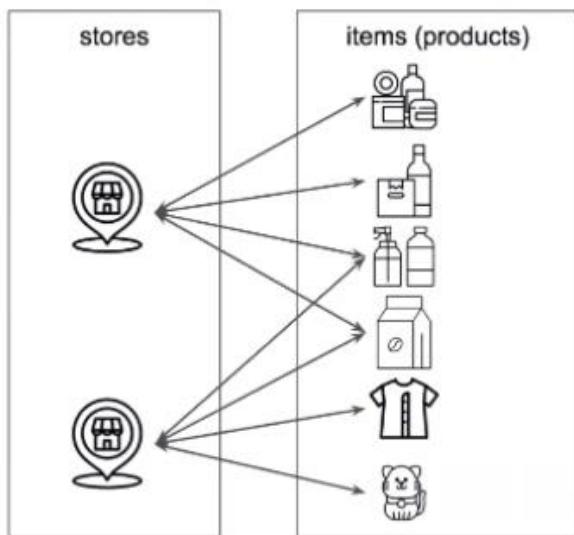
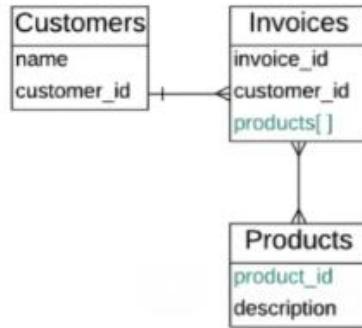
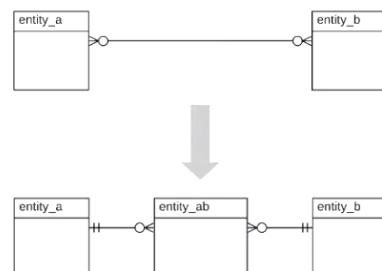


Common Relationships

- one-to-one (1-1)
- one-to-many (1-N)
- many-to-many (N-N)



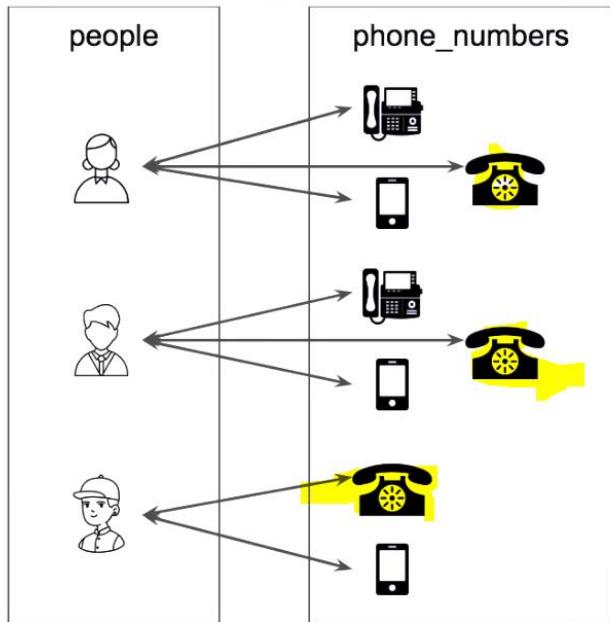
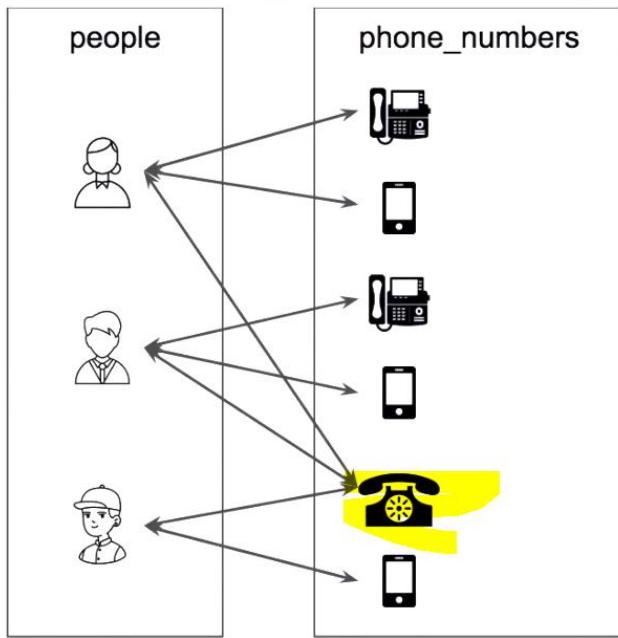
Many-to-Many => 2 x One-to-Many



actor_id	actor_name
actor1	Marlon Brando
actor2	Al Pacino
actor3	Robert DeNiro

movie_id	actor_id
movie1	actor1
movie1	actor2
movie2	actor2
movie2	actor3

movie_id	movie_title
movie1	The Godfather
movie2	The Godfather 2



Many-to-Many Representations

1. Embed

- a. array of subdocuments in the "many" side
- b. array of subdocuments in the other "many" side

Usually, only the most queried side is considered



2. Reference

- a. array of references in one "many" side
- b. array of references in the other "many" side

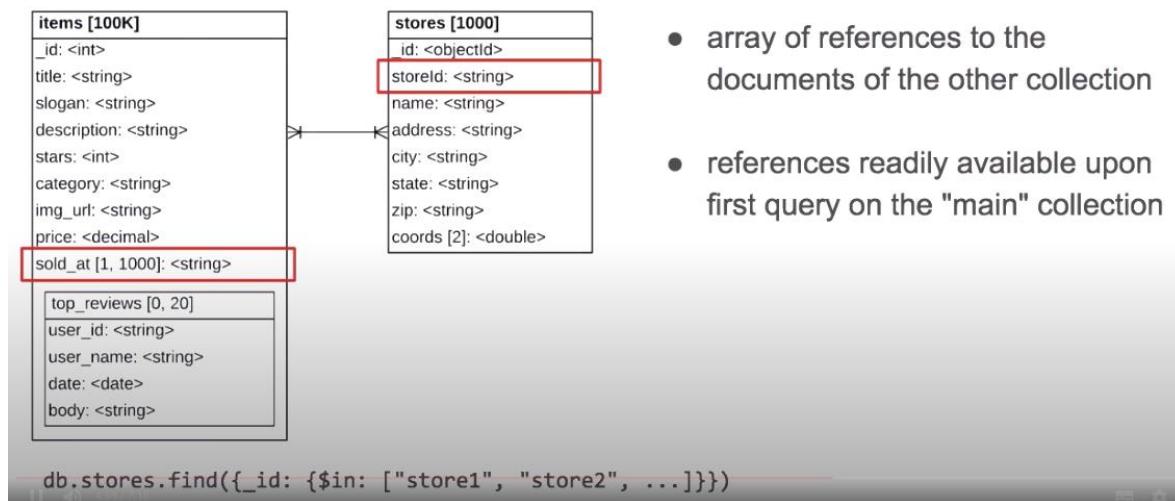
Many-to-Many: embed, in the main side

carts [10K]
_id: <objectId>
date: <date>
items: [0, 1, 1000]
qty: <int>
item
_id: <int>
title: <string>
stars: <int>
category: <string>
img_url: <string>
price: <decimal>

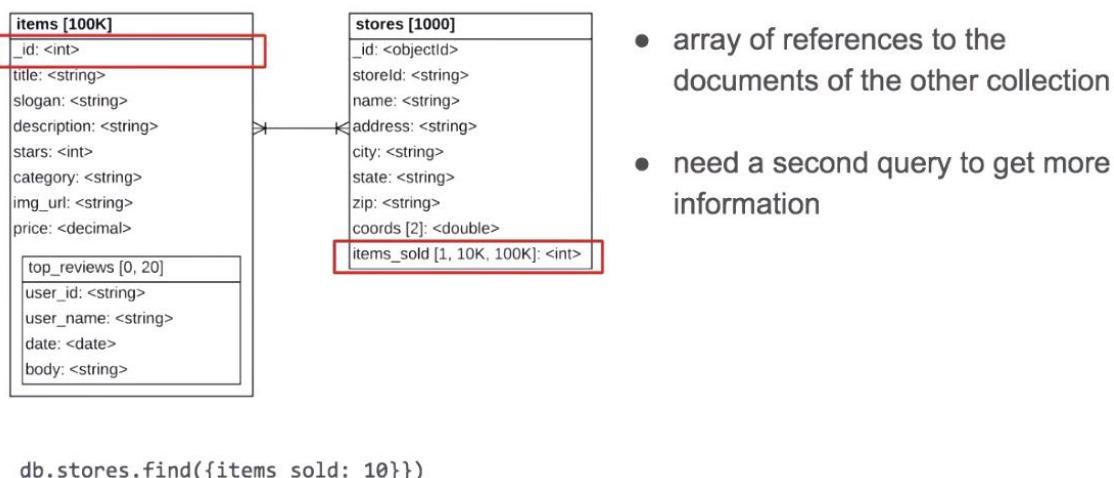
items [0, 100K]
_id: <int>
title: <string>
stars: <int>
category: <string>
img_url: <string>
price: <decimal>
quantity: <int>
...

- the documents from the less queried side are embedded
- results in duplication
- keep "source" for the embedded documents in another collection
- indexing is done on the array

Many-to-Many: reference, in the main side

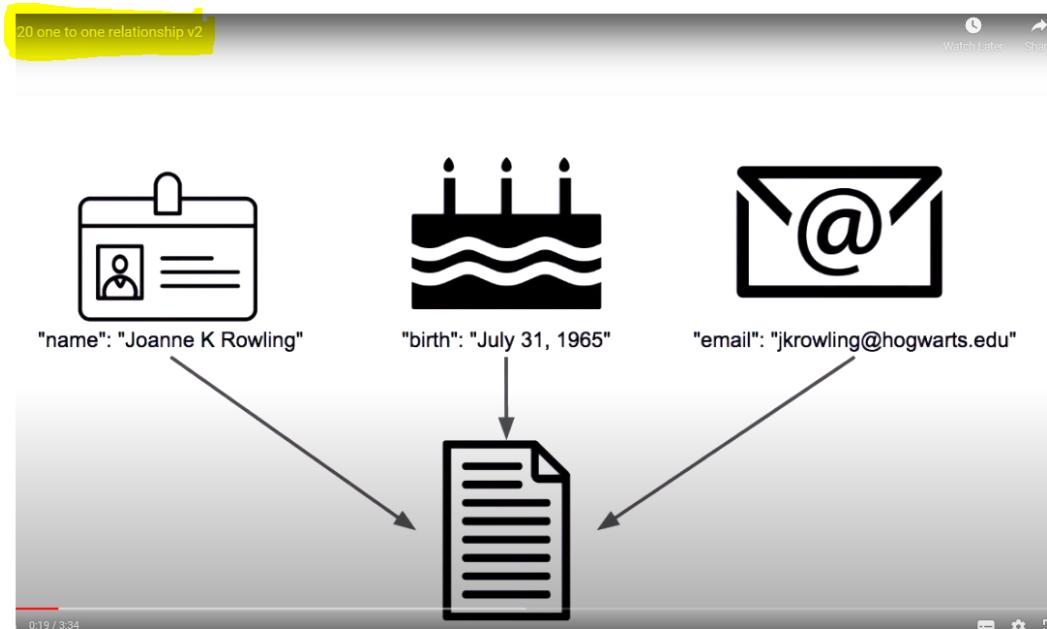


Many-to-Many: reference, in the secondary side



Recap for the Many-to-Many Relationships:

- Ensure it is a "many-to-many" relationship that should not be simplified.
- A "many-to-many" relationship can be replaced by two "one-to-many" relationships but does not have to with the document model
- Prefer embedding on the most queried side
- Prefer embedding for information that is primarily static over time and may profit from duplication.
- Prefer referencing over embedding to avoid managing duplication.



One-to-One Representations

1. Embed

- a. fields at same level
- b. grouping in sub-documents

2. Reference

- a. same identifier in both documents
- b. in the main "one" side
- c. in the secondary "one" side

One-to-One: embed, fields at the same level

- very similar to tabular databases

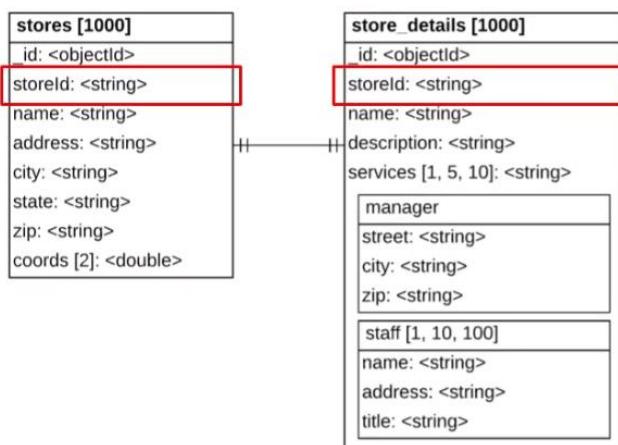
users [10M]
<code>_id: <objectId></code>
<code>name: <string></code>
<code>street: <string></code>
<code>city: <string></code>
<code>zip: <string></code>
<code>shipping_street: <string></code>
<code>shipping_city: <string></code>
<code>shipping_zip: <string></code>

One-to-One: embed, using subdocuments

users [10M]
_id: <objectId>
name: <string>
address
street: <string>
city: <string>
zip: <string>
shipping_address
street: <string>
city: <string>
zip: <string>

- preferred representation:
 - preserves simplicity
 - documents are clearer

One-to-One: reference

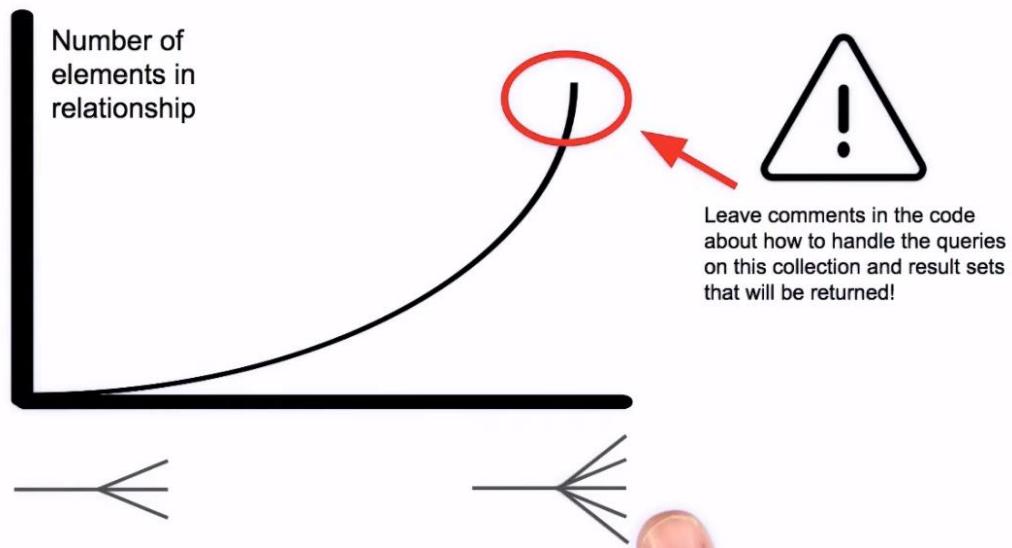
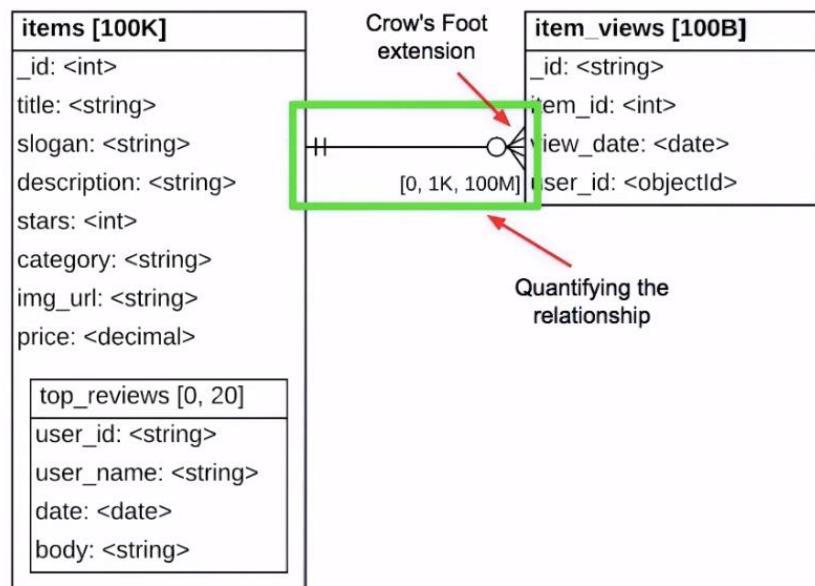


- add complexity
- possible performance improvements with:
 - smaller disk access
 - smaller amount of RAM needed

Recap for the One-to-One Relationships:

- Prefer embedding over referencing for simplicity
- Use subdocuments to organize the fields
- Use a reference for optimization purposes

Zillions !



One-to-Zillions Representations

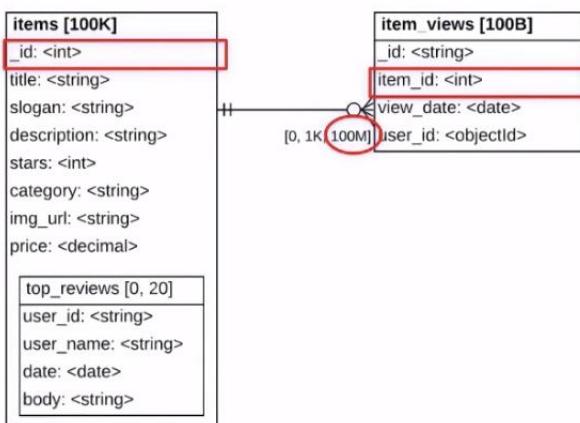
1. Embed

- a. in the "one" side
- b. in the "many" side

2. Reference

- a. in the "one" side
- b. in the "many/zillions" side

One-to-Zillions: reference, in the "zillions" side



- sole representation for a one-to-zillions relationship
- quantify the relationship to understand the maximum N value

Recap for the One-to-Zillions Relationships:

- It is a particular case of the one-to-many relationship.
- The only available representation is to reference the document on the "one" side of the relationship from the "zillion" side.
- Pay extra attention to queries and code that handle "zillions" of documents.

- Use computed pattern to avoid repetitive computations
- Structure similar fields with the attribute pattern
- Handle changes to your deployment with no downtime with a schema versioning pattern

Applying Patterns may lead to...

1. Duplication

duplicating data across documents

2. Data Staleness

accepting staleness in some pieces of data

3. Data Integrity Issues:

writing extra application side logic to ensure referential integrity

1. Duplication

Why?

- Result of embedding information in a given document for faster access

Concern?

- Challenge for correctness and consistency

1. Duplication - a. is the solution

```
{ // order
  date: "2009-02-14",
  status: "delivered",
  customer_id: "12345",
  items: [
    { item_id: "B000HC2LI0",
      quantity: 2 }
  ]
}

{ // customer
  _id: "12345",
  since: "2009-01-09",
  address: {
    street: "1600 Pennsylvania Ave",
    city: "Washington",
    state: "DC",
    country: "USA"
  }
}
```

1. Duplication - a. is the solution

```
{ // order
  date: "2009-02-14",
  status: "delivered",
  customer_id: "12345",
  items: [
    { item_id: "B000HC2LI0",
      quantity: 2 }
  ]
}

{ // customer
  _id: "12345",
  since: "2009-01-09",
  address: {
    street: "1600 Pennsylvania Ave",
    city: "Washington",
    state: "DC",
    country: "USA"
  }
}

{ // order
  date: "2009-02-14",
  status: "delivered",
  customer_id: "12345",
  items: [
    { item_id: "B000HC2LI0",
      quantity: 2 }
  ]
}

{ // customer
  _id: "12345",
  since: "2017-01-20",
  address: {
    street: "Hawaii Plantation Estate",
    city: "Paradise Point",
    state: "HI",
    country: "USA"
  }
}
```



1. Duplication - a. is the solution

```
{ // order
  date: "2009-02-14",
  status: "delivered",
  customer_id: "12345",
  items: [
    { item_id: "B000HC2LI0",
      quantity: 2 }
  ]
  shipping_address: {
    street: "1600 Pennsylvania Ave",
    city: "Washington",
    state: "DC",
    country: "USA"
  }
}
```



```
{ // order
  date: "2017-01-21",
  status: "delivered",
  customer_id: "12345",
  items: [
    { item_id: "B07CJWLLNX",
      quantity: 1 }
  ]
  shipping_address: {
    street: "Hawaii Plantation Estate",
    city: "Paradise Point",
    state: "HI",
    country: "USA"
  }
}
```



1. Duplication - b. has minimal effect

```
{ // movie
  _id: "tt0076759",
  title: "Star Wars - IV - A New Hope",
  cast: [
    "nm0000434",
    "nm0000148",
    ...
  ]
}

{ // actor
  _id: "nm0000434",
  name: "Mark Hamill",
  filmography: [
    { "tt0076759": "Luke Skywalker" },
    { "tt0080684": "Luke Skywalker" },
    ...
  ]
}
```

```
{ // movie 1
  title: "Star Wars - IV - A New Hope",
  cast: [
    { "Mark Hamill": "Luke Skywalker" },
    { "Harrison Ford": "Han Solo" },
    ...
  ]
}

{ // movie 2
  title: "Star Wars - V - The Empire Strikes Back",
  cast: [
    { "Mark Hamill": "Luke Skywalker" },
    { "Harrison Ford": "Han Solo" },
    ...
  ]
}
```



1. Duplication - c. should be handled

```
{ // movie
  _id: "tt0076759",
  title: "Star Wars - IV - A New Hope",
  gross_revenues: 775000000
}

{ // screenings
  date: "1977-05-30",
  revenues: 15554475,
}
...
{ // screenings
  date: "2019-05-30",
  revenues: 25000,
}
```



1. Duplication - c. should be handled

```
{ // movie           { // movie           { // movie
  _id: "tt0076759",   _id: "tt0076759",   _id: "tt0076759",
  title: "Star Wars - IV", title: "Star Wars - IV", title: "Star Wars - IV",
  gross_revenues: 775000000 gross_revenues: 775000000 gross_revenues: 775025000
}

{ // screenings     { // screenings     { // screenings
  date: "1977-05-30", date: "1977-05-30", date: "1977-05-30",
  revenues: 15554475,  revenues: 15554475,  revenues: 15554475,
}

...
{
  date: "2019-05-30", date: "2019-05-30", date: "2019-05-30",
  revenues: 5000,      revenues: 25000,      revenues: 25000,
}

{
  date: "2019-06-17", date: "2019-06-17", date: "2019-06-17",
  revenues: 25000,      revenues: 25000,      revenues: 25000,
}
```



m320 handling duplication part 3

2. Staleness

Why?

- New events come along at such a rate that updating some data constantly causes performance issues.

Concern?

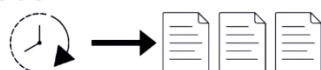
- Data quality and reliability.

Long analytic queries?

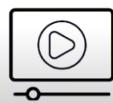
Queries on secondary nodes?

2. Resolve Staleness

Batch Updates



Find the changes through a
Change Stream





Using the Attribute Pattern

```

{
  "manufacturer": "China",
  "brand": "MongoDB",
  "sub_brand": "University",
  "price": 0.00,
  ...
  "color": "black",
  "size": "100 x 70 x 10 mm"
  ...
  "input": "5V/1300 mA"
  "output": "5V/1A"
  "capacity": "4200 mAh"
}

{
  "manufacturer": "China",
  "brand": "MongoDB",
  "sub_brand": "University",
  "price": 0.00,
  ...
  "color": "black",
  "size": "100 x 70 x 10 mm"
  ...
  "add_specs": [
    { "k": "input", "v": "5V/1300 mA" },
    { "k": "output", "v": "5V/1A" },
    { "k": "capacity", "v": 4200, "u": "mAh" }
  ]
}

```

Fields that share Common Characteristics

```

{
  "title": "Dunkirk",
  ...
  "release_USA": "2017/07/23",
  "release_Mexico": "2017/08/01",
  "release_France": "2017/08/01",
  "release_Festival_San_Jose":
    "2017/07/22"
}

{
  "title": "Dunkirk",
  ...
  "releases": [ {
    { "k": "release_USA",
      "v": "2017/07/23" },
    { "k": "release_Mexico",
      "v": "2017/08/01" },
    { "k": "release_France",
      "v": "2017/08/01" },
    { "k": "release_Festival_San_Jose",
      "v": "2017/07/22" }
  ]
}

db.movies.find({ "releases.v": { $gte: "2017/07", $lt: "2017/08" } })

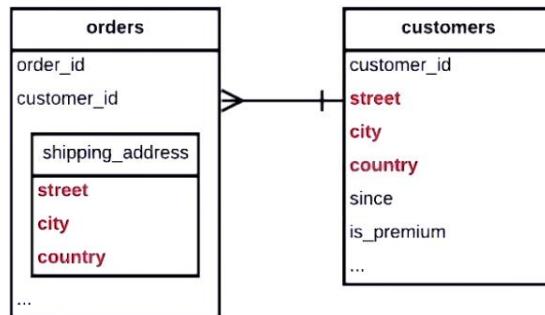
```

Attribute Pattern

Problem	Solution
<ul style="list-style-type: none">• Lots of similar fields• Want to search across many fields at once• Fields present in only a small subset of documents	<ul style="list-style-type: none">• Break the <i>field/value</i> into a sub-document with:<ul style="list-style-type: none">◦ fieldA: <i>field</i>◦ fieldB: <i>value</i>• Example:<ul style="list-style-type: none">◦ { "color": "blue", "size": "large" }◦ { [{ "k": "color", "v": "blue" }, { "k": "size", "v": "large" }] }
Use Cases Examples	Benefits and Trade-Offs
<ul style="list-style-type: none">• Characteristics of a product• Set of fields all having same value type<ul style="list-style-type: none">◦ List of dates	<ul style="list-style-type: none">✓ Easier to index✓ Allow for non-deterministic field names✓ Ability to qualify the relationship of the original field and value

Extended Reference for Many-to-One Relationships

- embed the "One" side, of a "One-to-Many" relationship, into the "Many" side
- only the part that we need to join often

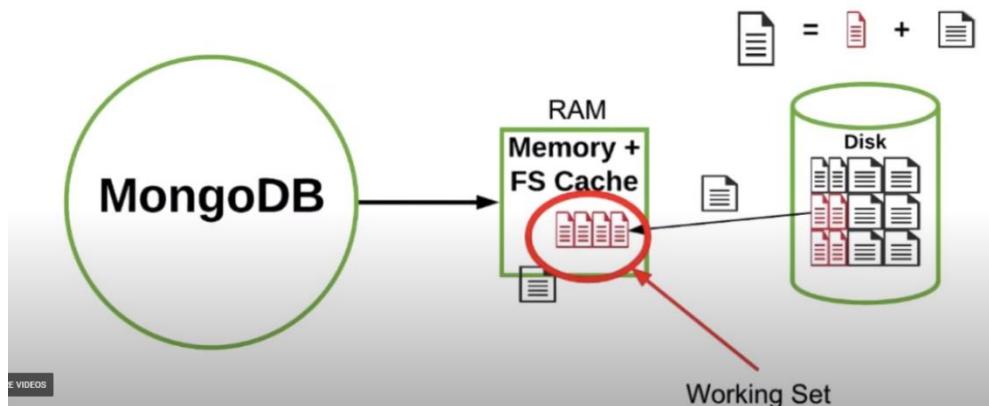


Extended Reference Pattern

Extended Reference Pattern

Problem	Solution
<ul style="list-style-type: none">• Too many repetitive joins	<ul style="list-style-type: none">• Identify fields on the lookup side• Bring those fields into the main object
Use Cases Examples	Benefits and Trade-Offs
<ul style="list-style-type: none">• Catalog• Mobile Applications• Real-Time Analytics	<ul style="list-style-type: none">✓ Faster reads✓ Reduce number of joins and lookups✗ May introduce lots of duplication if extended reference contains fields that mutate a lot

Reduce the size of the Working Set



Subset Pattern

Problem

- Working set is too big
- Lot of pages are evicted from memory
- A large part of documents is rarely needed

Solution

- Split the collection in 2 collections
 - Most used part of documents
 - Less used part of documents
- Duplicate part of a 1-N or N-N relationship that is often used in the most used side

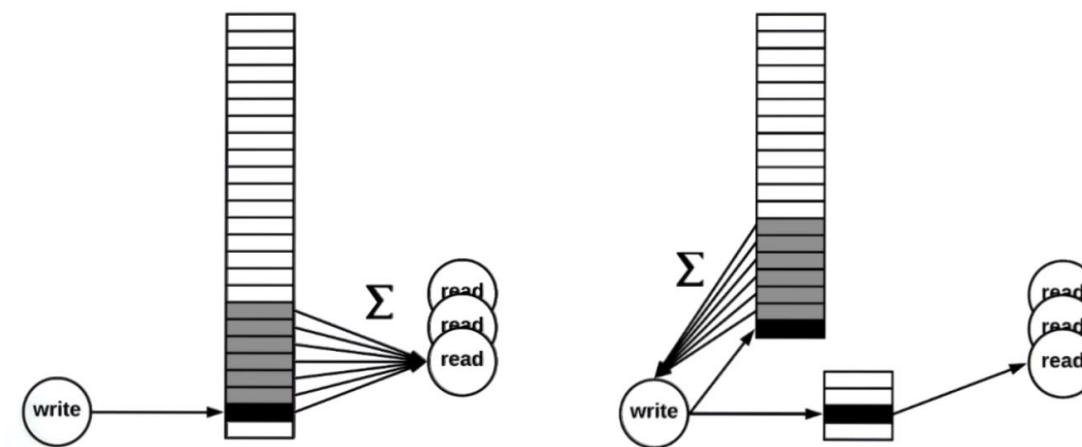
Use Cases Examples

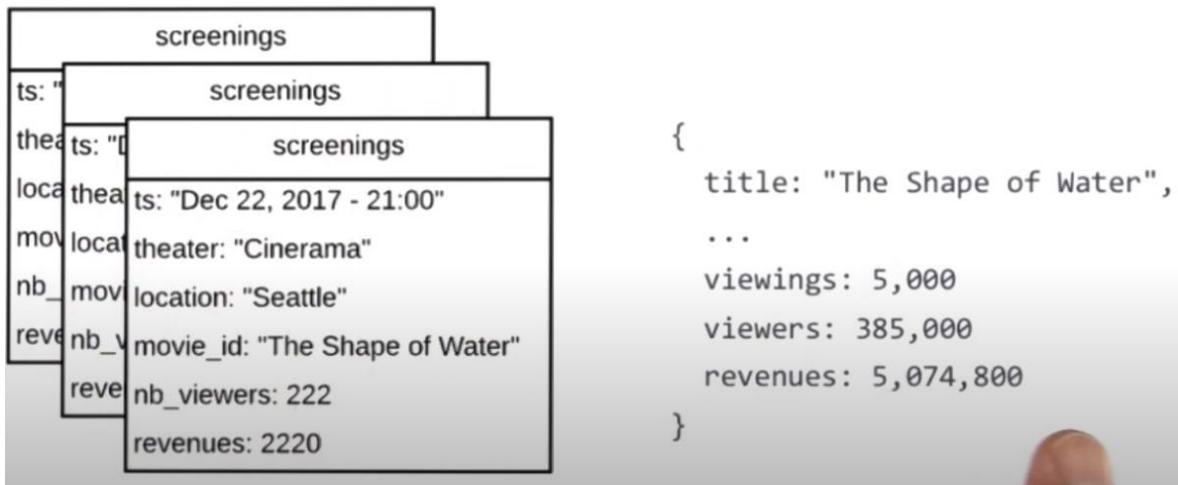
- List of reviews for a product
- List of comments on an article
- List of actors in a movie

Benefits and Trade-Offs

- ✓ Smaller working set, as often used documents are smaller
- ✓ Shorter disk access for bringing in additional documents from the most used collection
- ✗ More round trips to the server
- ✗ A little more space used on disk

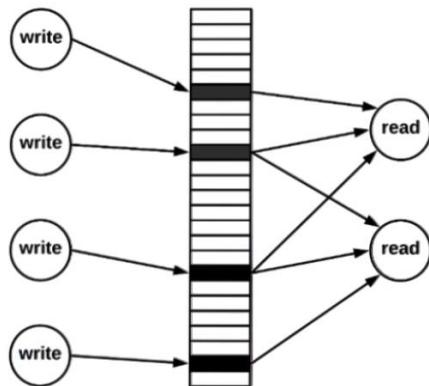
Mathematical Operations



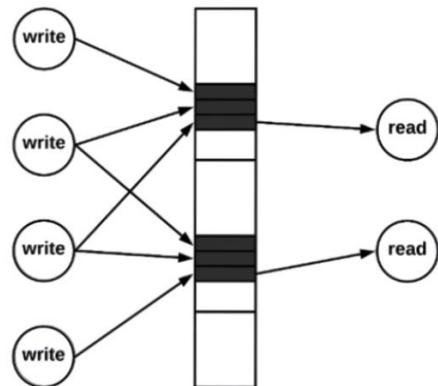


Fan Out Operations

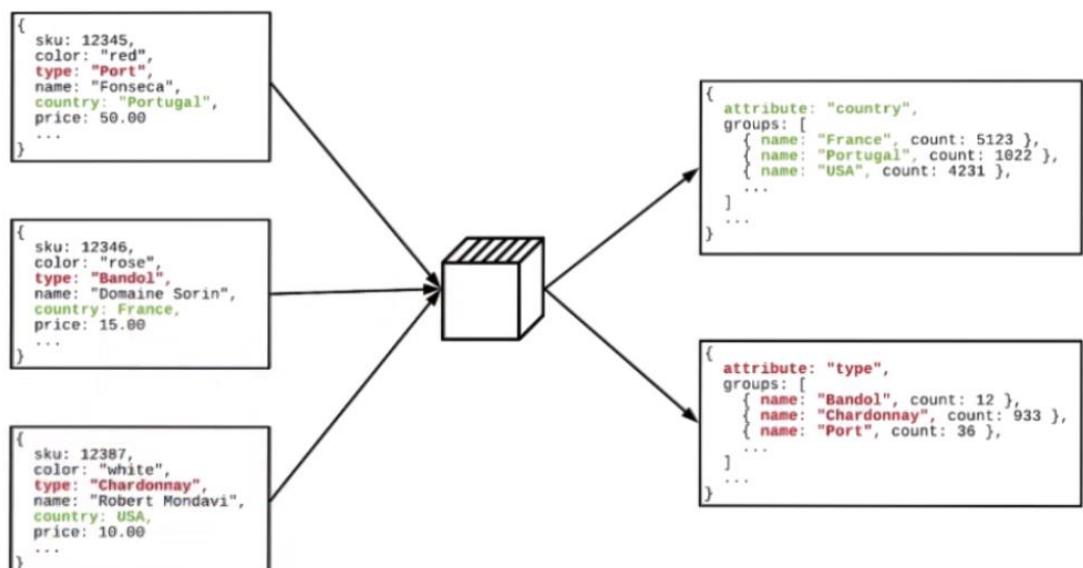
Fan Out on Reads



Fan Out on Writes



Example of Roll Up



Computed Pattern

Problem

- Costly computation or manipulation of data
- Executed frequently on the same data, producing the same result

Solution

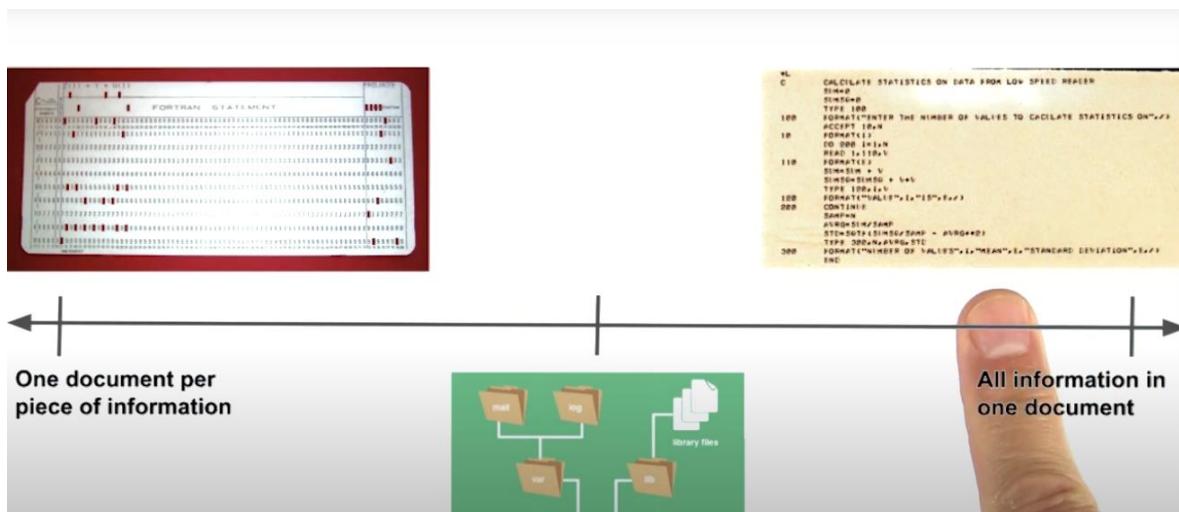
- Perform the operation and store the result in the appropriate document and collection
- If need to redo the operations, keep the source of them

Use Cases Examples

- Internet Of Things (IOT)
- Event Sourcing
- Time Series Data
- Frequent Aggregation Framework queries

Benefits and Trade-Offs

- ✓ Read queries are faster
- ✓ Saving on resources like CPU and Disk
- ✗ May be difficult to identify the need
- ✗ Avoid applying or overusing it unless needed



Just the Right Amount

One document per device per day

```
{  
  "device_id": 000123456,  
  "type": "2A",  
  "date": ISODate("2018-03-02"),  
  "temp": [ [ 20.0, 20.1, 20.2, ... ],  
            [ 22.1, 22.1, 22.0, ... ],  
            ...  
          ]  
  
}  
  
{  
  "device_id": 000123456,  
  "type": "2A",  
  "date": ISODate("2018-03-03"),  
  "temp": [ [ 20.1, 20.2, 20.3, ... ],  
            [ 22.4, 22.4, 22.3, ... ],  
            ...  
          ]  
  
}
```

MORE VIDEOS

One document per device per hour

```
{  
  "device_id": 000123456,  
  "type": "2A",  
  "date": ISODate("2018-03-02T13"),  
  "temp": { 1: 20.0, 2: 20.1, 3: 20.2, ... }  
}  
  
{  
  "device_id": 000123456,  
  "type": "2A",  
  "date": ISODate("2018-03-02T14"),  
  "temp": { 1: 22.1, 2: 22.1, 3: 22.0, ... }  
}
```



Collaboration Platform Example

- Many channels
- Design alternatives
 - One document per message
 - One document for all messages in a channel

```
{  
  "channel_id": "13579246",  
  "name": "mongodb_tech",  
  "date": ISODate("2018-03-02T03:14:16"),  
  "message": "<Daniel joined the group>"  
}  
  
{  
  "channel_id": "13579246",  
  "name": "mongodb_tech",  
  "messages": [  
    "<Channel created>",  
    "<Nathan joined the group>",  
    "<Daniel joined the group>",  
    "Nathan: Welcome Daniel",  
    "Daniel: Thanks!"  
    ...  
  ]  
}
```



Collaboration Platform Example

- Many channels
- Design alternatives
 - One document per message
 - One document for all messages in a channel
 - One document per channel, per day

```
{  
  "channel_id": "13579246",  
  "name": "mongodb_tech",  
  "date": ISODate("2018-03-02"),  
  "messages": [  
    "<Daniel joined the group>",  
    "Nathan: Welcome Daniel",  
    "Daniel: Thanks!"  
    ...  
  ]  
}
```

Column Oriented Data

```
{  
  "device_id": 000123456,  
  "location": [ 85.1, 17.2 ]  
  "date": ISODate("2018-03-02T13"),  
  "temp": [ 20.0, 20.1, 20.2, ... ],  
  "pressure": [ 1.01, 1.02, 1.01, ... ],  
  "light": [ 8500, 8520, 8525, ... ],  
  ...  
}  
  
db.iot.aggregate([{$project:{avg_temp:{$avg:  
    "$temp" }}}])  
  
{  
  "device_id": 000123456,  
  "date": ISODate("2018-03-02T13"),  
  "temp": [ 20.0, 20.1, 20.2, ... ]  
}  
  
{  
  "device_id": 000123456,  
  "date": ISODate("2018-03-02T13"),  
  "pressure": [ 1.01, 1.02, 1.01, ... ]  
}  
  
{  
  "device_id": 000123456,  
  "date": ISODate("2018-03-02T13"),  
  "light": [ 8500, 8520, 8525, ... ]  
}  
  
db.iot.aggregate([{$match:{temp:$exists:1}},  
  {$project:{avg_temp:{$avg:"$temp" }}}])
```

Gotchas with Buckets

- Random insertions or deletions in buckets
- Difficult to sort across buckets
- Ad hoc queries may be more complex, again across buckets
- Works best when the "complexity" is hidden through the application code

Bucket Pattern

Problem	Solution
<ul style="list-style-type: none">• Avoiding too many documents, or too big documents• A 1-to-Many relationship that can't be embedded	<ul style="list-style-type: none">• Define the optimal amount of information to group together• Create arrays to store the information in the main object• It is basically an embedded 1-to-Many relationship, where you get N documents, each having an average of Many/N sub documents.

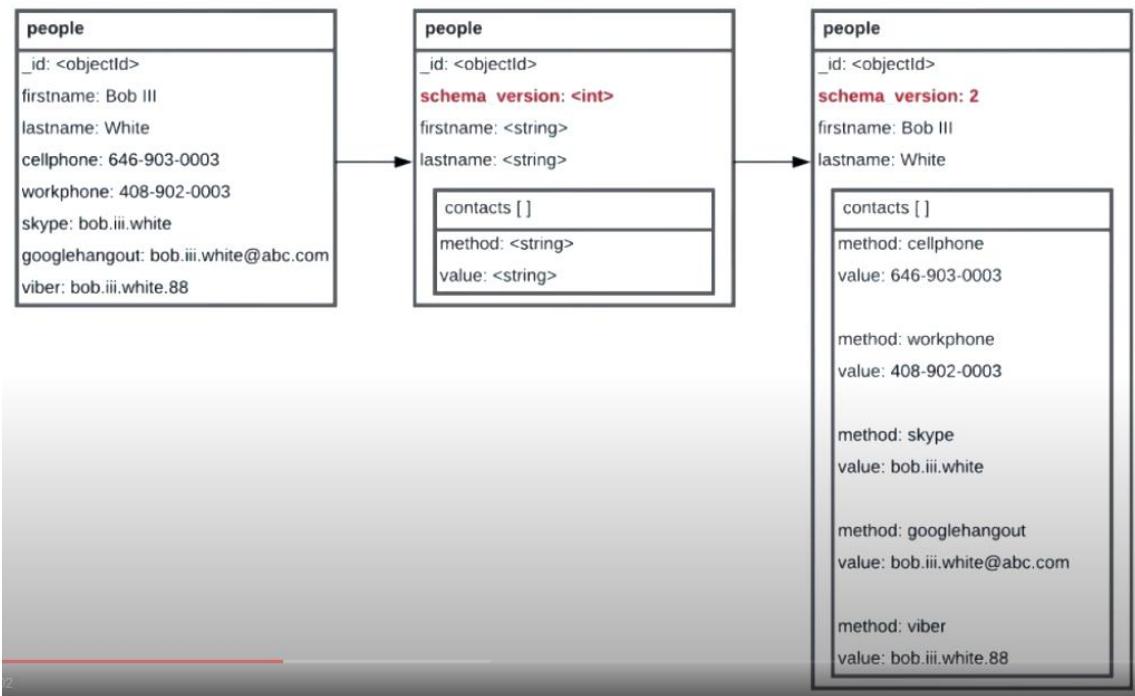
Use Cases Examples	Benefits and Trade-Offs
<ul style="list-style-type: none">• Internet Of Things• Data Warehouse• Lots of information associated to one object	<ul style="list-style-type: none">✓ Good balance between number of data access and size of data returned✓ Makes data more manageable✓ Easy to prune data✗ Can lead to poor query results if not designed correctly✗ Less friendly to BI Tools

m320 Schema Versioning

people	people	people
_id: <objectId> firstname: Bob Sr. lastname: White homephone: 408-901-0001 workphone: 408-902-0001	_id: <objectId> firstname: Bob Jr. lastname: White homephone: 408-901-0002 workphone: 408-902-0002 cellphone: 646-903-0002	_id: <objectId> firstname: Bob III lastname: White cellphone: 646-903-0003 workphone: 408-902-0003 skype: bob.iii.white googlehangout: bob.iii.white@abc.com viber: bob.iii.white.88

VIDEOS

2.01 / 7.02

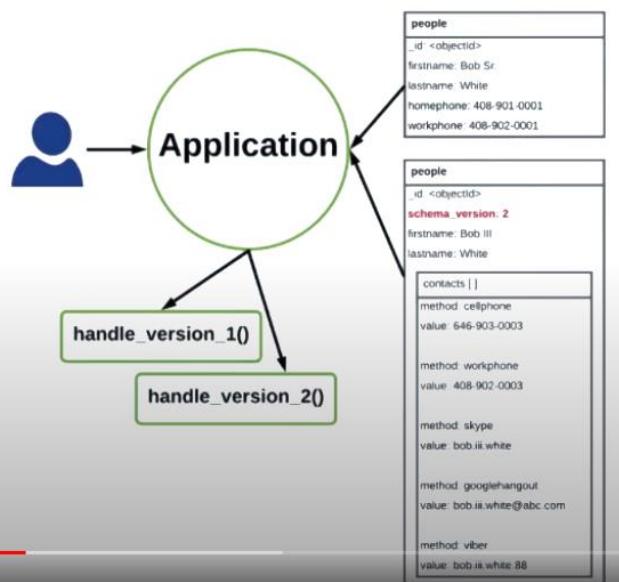


Application Lifecycle

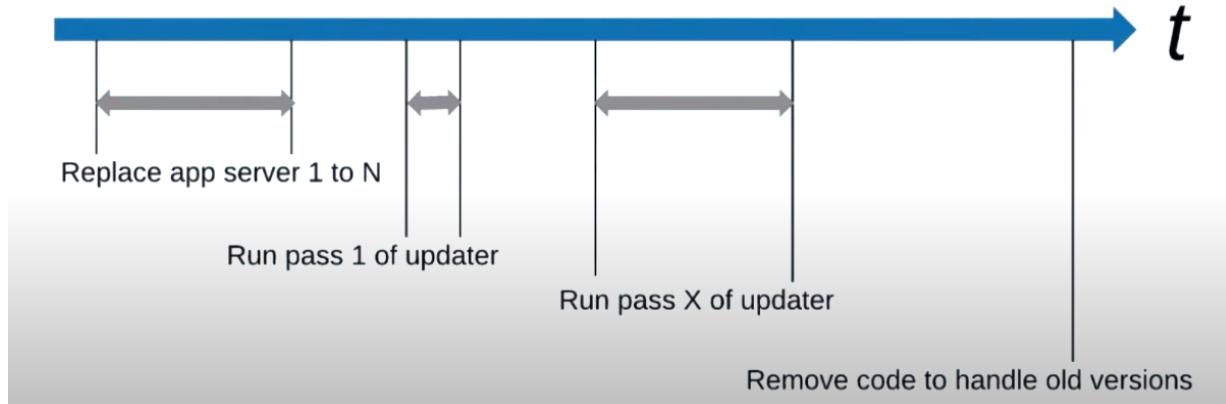
- Modify Application
 - Can read/process all versions of documents
 - Have different handler per version
 - Reshape the document before processing it
- Update all Application servers
 - Install updated application
 - Remove old processes
- Once migration completed
 - remove the code to process old versions

VIDEOS

2:56 / 7:02



Timeline of the Migration



Schema Versioning Pattern

Problem	Solution
<ul style="list-style-type: none">• Avoid downtime while doing schema upgrades• Upgrading all documents can take hours, days or even weeks when dealing with big data.• Don't want to update all documents	<ul style="list-style-type: none">• Each document gets a "schema_version" field• Application can handle all versions• Choose your strategy to migrate the documents

Use Cases Examples

- Every application that use a database, deployed in production and heavily used.
- System with a lot of legacy data

Benefits and Trade-Offs

- ✓ No downtime needed
- ✓ Feel in control of the migration
- ✓ Less future technical debt
- ✗ May need 2 indexes for same field while in migration period

m320 tree patterns v3

Who are the ancestors of node X?

Who reports to Y ?

Find all nodes that are under Z ?

Change all categories under N to under P

```
{  
    "_id" : 7,  
    "title" : "Brown Tumbler",  
    "slogan" : "Bring your coffee to go",  
    "description" : "The MongoDB Insulated Travel Tumbler is smartly designed to maintain temperatures ",  
    "stars" : 0,  
    "category" : "Kitchen",  
    "img_url" : "/img/products/brown-tumbler.jpg",  
    "price" : NumberDecimal("9.00"),  
    "sold_at" : [  
        ObjectId("55aea9699f876cd5bcb77fcc"),  
        ObjectId("55aea9699f876cd5bcb77fdb"),  
        ...  
    ],  
    "top_reviews" : [  
        {  
            "body" : "Customer finish player that.",  
            "user_id" : ObjectId("5cc8beb8b2c78b148aededf6"),  
            "user_name" : "Cody Allen",  
            "date" : "1983-04-18"  
        },  
        ...  
    ]  
}
```

Model Tree Structures

Parent References

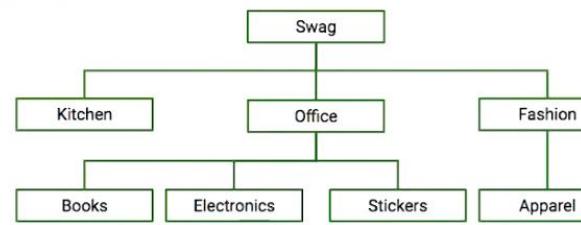
Child References

Array of Ancestors

Materialized Paths

Parent References

```
{  
  name: "Office",  
  parent: "Swag",  
  ...  
}
```

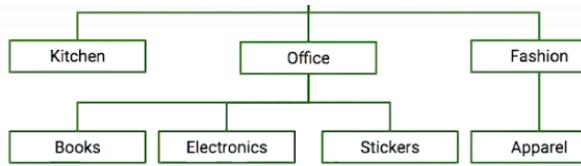


Who are the ancestors of node X?

```
// all ancestors  
db.categories.aggregate([  
  { $graphLookup: {  
    from: 'categories',  
    startWith: '$name',  
    connectFromField: 'parent',  
    connectToField: 'name',  
    as: 'ancestors',  
  }}])
```

Parent References

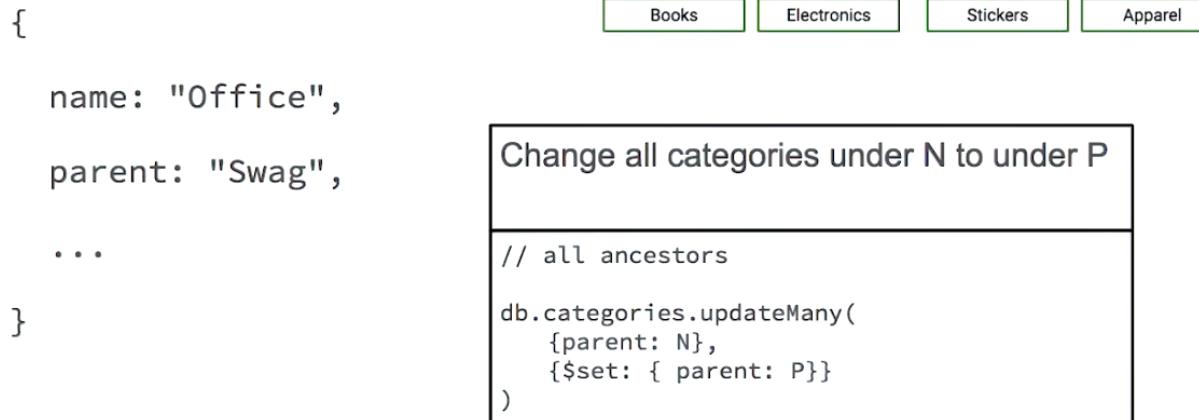
```
{  
  name: "Office",  
  parent: "Swag",  
  ...  
}
```



Who reports to Y?

```
// immediate ancestor  
db.categories.find({parent: Y})
```

Parent References



Parent References

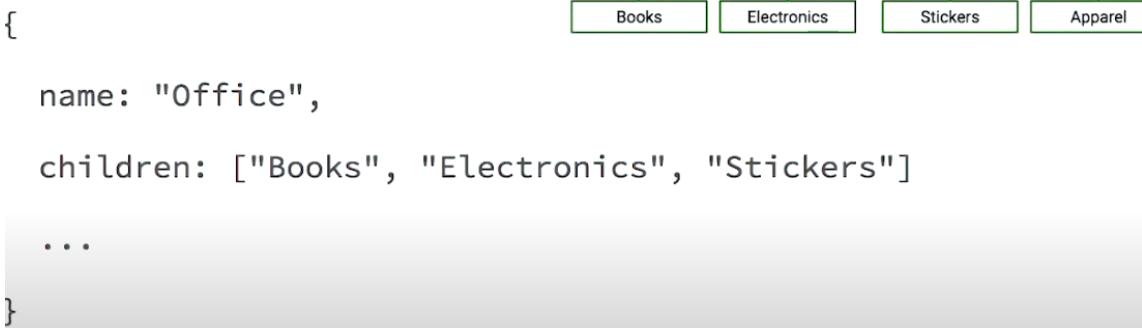
Who are the ancestors of node X? !

Who reports to Y? ✓

Find all nodes that are under Z ? !

Change all categories under N to under P ✓

Child References



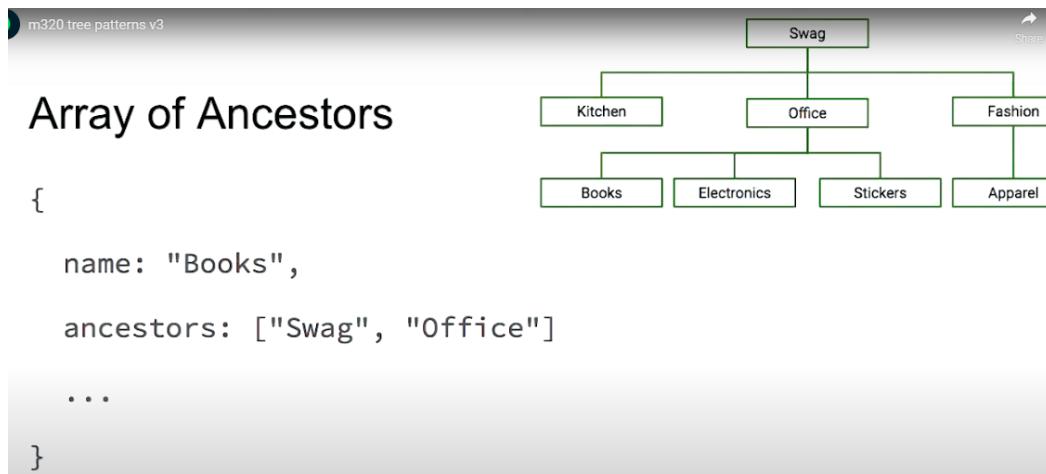
Child References

Who are the ancestors of node X? !

Who reports to Y ? !

Find all nodes that are under Z ? ✓

Change all categories under N to under P !



Array of Ancestors

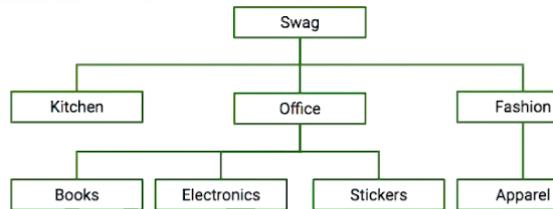
Who are the ancestors of node X? ✓

Who reports to Y ? ✓

Find all nodes that are under Z ? !

Change all categories under N to under P !

Materialized Paths



```
{
```

```
    name: "Books",  
    ancestors: ".Swag.Office"
```

```
    ...
```

```
}
```

```
// immediate ancestor of Y  
db.categories.find({ancestors: /\.\Y$/})  
  
// if descends from X and Z  
db.categories.find({ancestors: /^\.\X.*\Y/i})
```

Materialized Paths

Who are the ancestors of node X? ✓

Who reports to Y ? !

Find all nodes that are under Z ? !

Change all categories under N to under P !

MongoMart Solution: Ancestor Array + Parent Reference

```
{  
    "_id" : 8,  
    "name" : "Umbrellas",  
    "parent" : "Fashion",  
    "ancestors" : [ "Swag", "Fashion" ]  
}
```

MongoMart Solution: Ancestor Array + Parent Reference

```
{  
  "_id" : 8,  
  "name" : "Umbrellas",  
  "parent" : "Fashion",  
  "ancestors" : [ "Swag", "Fashion" ]  
}
```

```
// Navigate up on the tree  
db.categories.find({parent: X})  
  
// Find all nodes under a given category  
db.categories.find({ancestors: Y })
```

Tree Patterns



Problem

- Representation of hierarchical structured data
- Different access patterns to navigate the tree
- Provide optimized model for common operations

Solution

- Different patterns
 - Child Reference
 - Parent Reference
 - Array of Ancestors
 - Materialized Paths

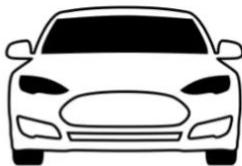
Use Cases Examples

- Org Charts
- Product Categories

Benefits and Trade-Offs

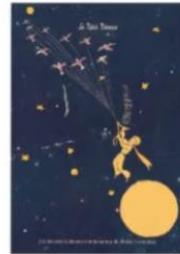
- Child Reference
- ✓ Easy to navigate to children nodes or tree descending access patterns
- Parent Reference
- ✓ Immediate parent node discovery and tree updates
- Array of Ancestors
- ✓ Navigate upwards on the ancestors path
- Materialized Path
- ✓ Makes use of regular expressions to find nodes in the tree

Vehicles



```
{ "vehicle_type": "car", ... "owner": "Yulia", "taxes": "200", "wheels": 4 } { "vehicle_type": "truck", ... "owner": "Daniel", "taxes": "800", "wheels": 10, "axles": 3. } { "vehicle_type": "boat", ... "owner": "Norberto", "taxes": "2000" }
```

Products



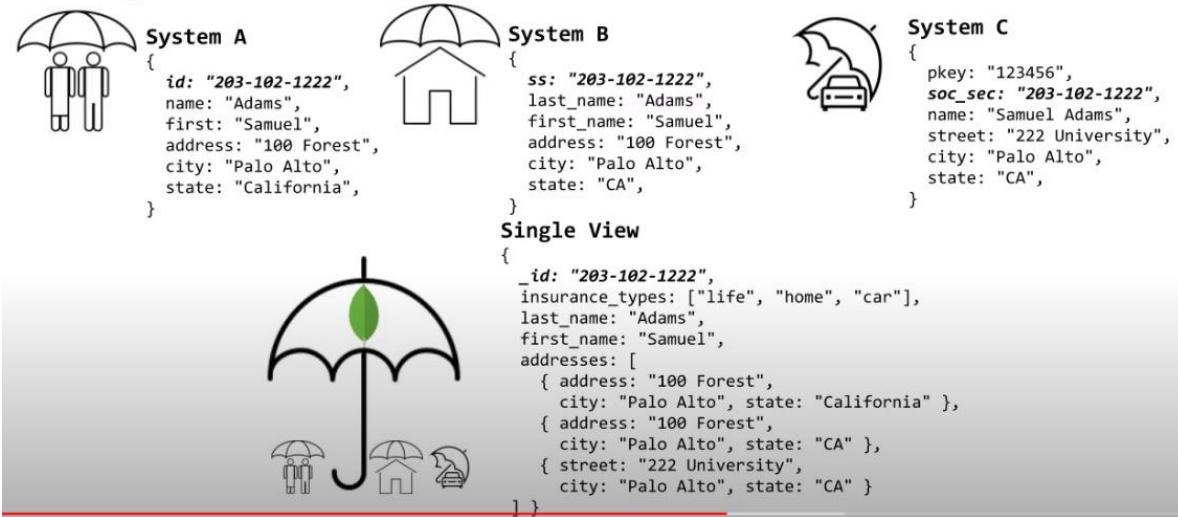
```
{ "product_type": "shirt", "color": "blue", "size": "large", "price": "100.00" } { "product_type": "book", "title": "Le Petit Prince", "size": "20cm x 15cm x 1cm", "price": "10.00" }
```

Polymorphic Subdocuments

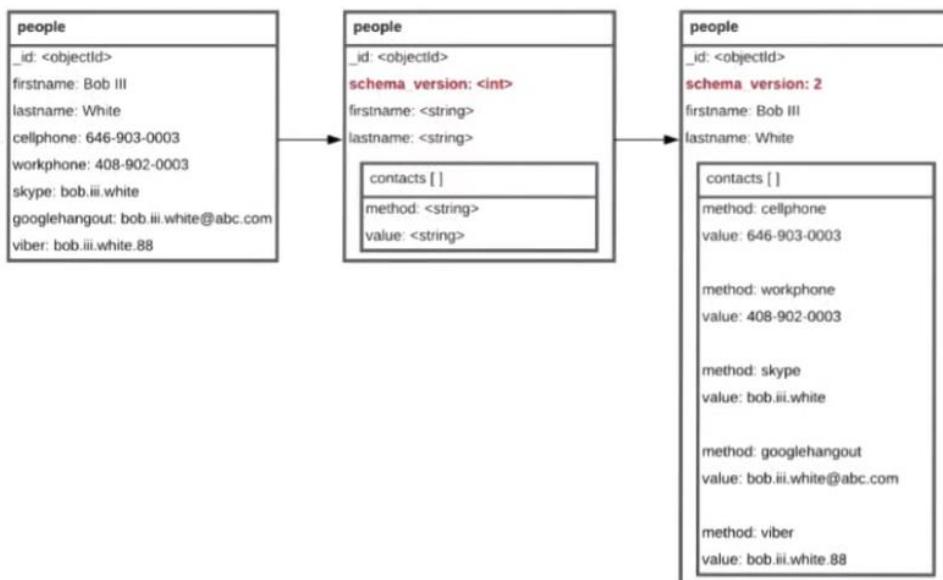
```
People
{
  "first_name": "Norberto",
  ...
  "addresses": [ {
    "street": "725 5th Ave"
    "city": "New York",
    "state": "NY",
    "zip_code": "10022",
    "country": "USA"
  },
  ...
  ]
}
```

```
People
{
  "first_name": "Norberto",
  ...
  "addresses": [ {
    "street": "725 5th Ave"
    "city": "New York",
    "state": "NY",
    "zip_code": "10022",
    "country": "USA"
  },
  {
    "street": "Palacio de Belem",
    "city": "Lisbon",
    "postal_code": "1300-004",
    "country": "Portugal"
  }
}
```

Single View Example



Polymorphism in the Schema Versioning Pattern



Polymorphic Pattern

Problem

- Objects more similar than different
- Want to keep objects in same collection

Solution

- Field tracks the type of document or sub-document
- Application has different code paths per document type, or has subclasses

Use Cases Examples

- Single View
- Product Catalog
- Content Management

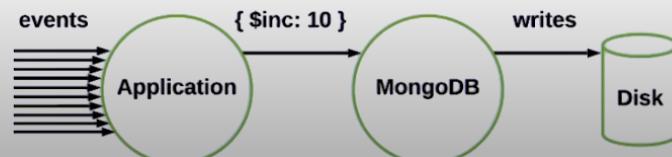
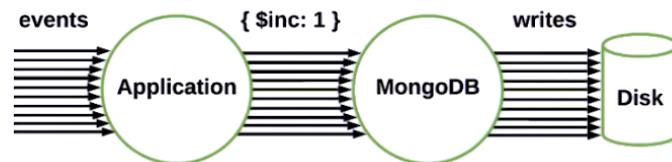
Benefits and Trade-Offs

- ✓ Easier to implement
- ✓ Allow to query across a single collection

Additional Patterns

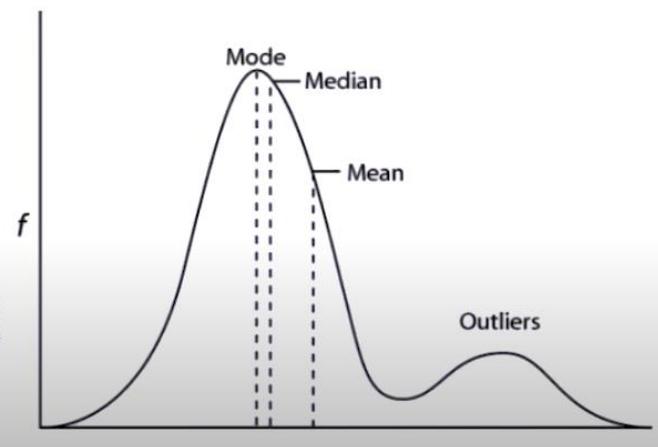
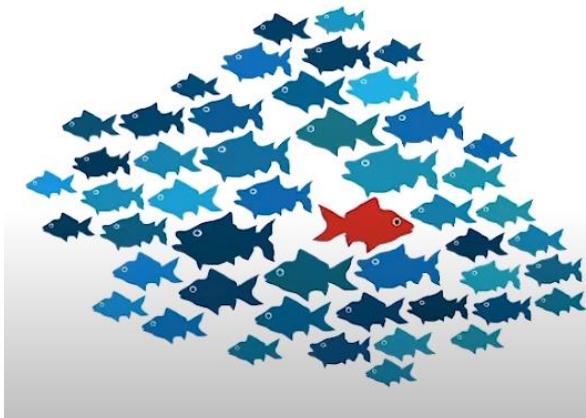
- Approximation Pattern
- Outlier Pattern

Approximation Pattern



Approximation Pattern	
Problem	Solution
<ul style="list-style-type: none">• Data is expensive to calculate• It does not matter if the number is not precise	<ul style="list-style-type: none">• Fewer, but writes with higher payload
Use Cases Examples	Benefits and Trade-Offs
<ul style="list-style-type: none">• Web page counters• Any counters with tolerance to imprecision• Metric statistics	<ul style="list-style-type: none">✓ Less writes✓ Less contention on documents✓ Statistically valid numbers✗ Not exact numbers✗ Must be implemented in the application

Outlier Pattern



Movies Site Solution

Collection: movie_extras

```
{  
  "_id": 1980001234,  
  "title": "Gandhi",  
  "has_overflow_extras": true,  
  "extras": [  
    { "first_name": "Angshuman",  
      "Last_name": "Bagchi" },  
    { "first_name": "Rupa",  
      "last_name": "Narayanan" },  
    ...  
  ],  
  number_extras: 1000  
}
```

```
{  
  "_id": 1980001234001,  
  "movie_id": 1980001234,  
  "title": "Gandhi",  
  "is_overflow_extra": true,  
  "extras": [  
    { "first_name": "Sachin",  
      "Last_name": "Tendulkar" },  
    { "first_name": "Sourav",  
      "last_name": "Ganguly" },  
    ...  
  ],  
  number_extras: 820  
}
```

Outlier Pattern

Problem	Solution
<ul style="list-style-type: none">• Few documents would drive the solution• Impact would be negative on the majority of queries	<ul style="list-style-type: none">• Implementation that works for majority• Field identifies outliers as exception• Outliers are handled differently on the application side
Use Cases Examples	Benefits and Trade-Offs
<ul style="list-style-type: none">• Social Networks• Popularity	<ul style="list-style-type: none">✓ Optimized solution for most use cases✗ Differences handled application side✗ Difficult for aggregation of ad hoc queries

Recap:

These are some other notable patterns

- Approximation Pattern
 - avoiding performing an operation too often
- Outlier Pattern
 - keeping the focus on the most frequent use cases

